

FastKLEE: Faster Symbolic Execution via Reducing Redundant Bound Checking of Type-Safe Pointers

ESEC/FSE 2022 Tool Demonstration

Haoxin Tu, Lingxiao Jiang, Xuhua Ding (Singapore Management University)
He Jiang (Dalian University of Technology)



Outlines

- **Background**
- **Motivation**
- **Solution**
 - FastKLEE
- **Preliminary Evaluation**
- **Conclusion**

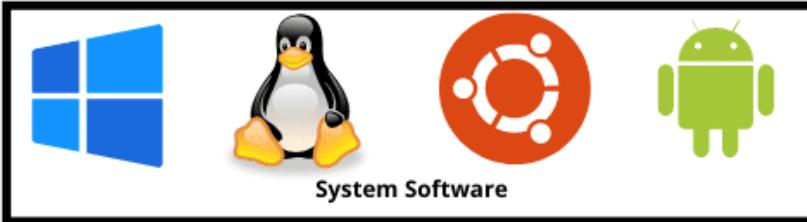
Background



School of
Information Systems

Background

Examples of System & Application Software



Background

Examples of System & Application Software



- Q: How to improve the quality of software?

Background

Examples of System & Application Software



- Q: How to improve the quality of software?
- A: Software testing

Background

Examples of System & Application Software



- Q: How to improve the quality of software?
- A: Software testing



(Symbolic Execution)

Background

Examples of System & Application Software



- Q: How to improve the quality of software?
- A: Software testing



(Symbolic Execution)



1

A General Workflow of Symbolic Execution (SE)

Background

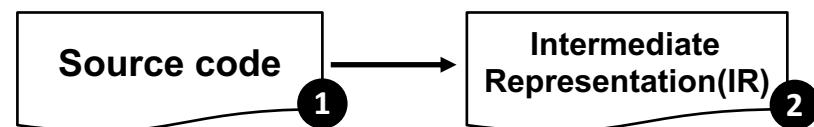
Examples of System & Application Software



- Q: How to improve the quality of software?
- A: Software testing



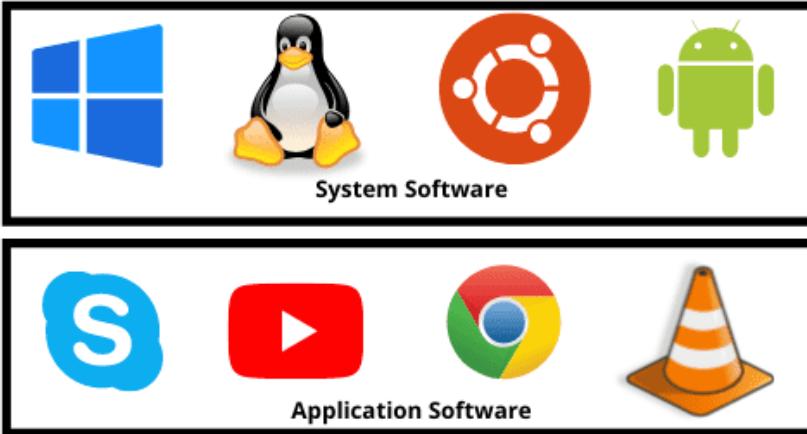
(Symbolic Execution)



A General Workflow of Symbolic Execution (SE)

Background

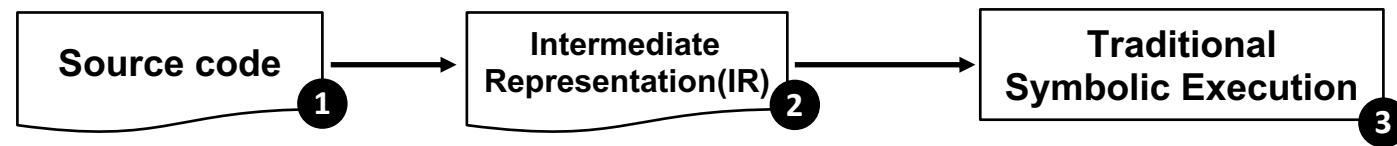
Examples of System & Application Software



- Q: How to improve the quality of software?
- A: Software testing



(Symbolic Execution)



A General Workflow of Symbolic Execution (SE)

Background

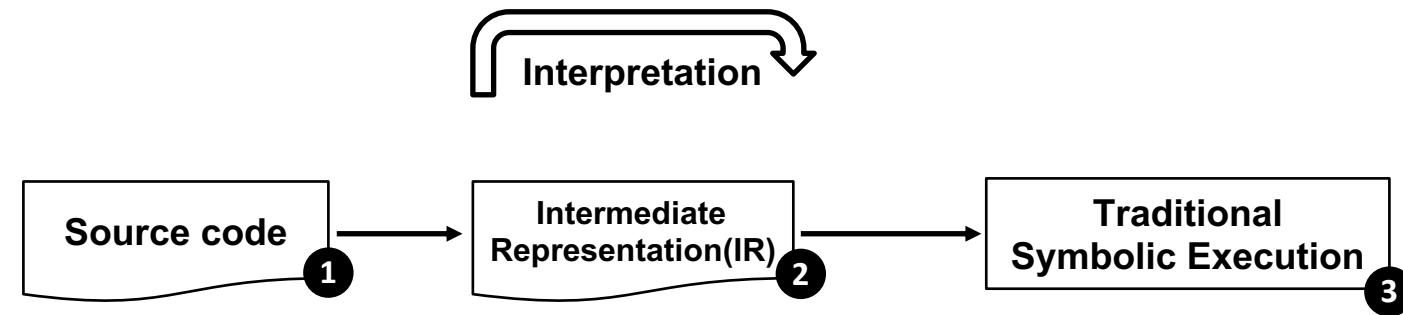
Examples of System & Application Software



- Q: How to improve the quality of software?
- A: Software testing



(Symbolic Execution)



A General Workflow of Symbolic Execution (SE)

Background

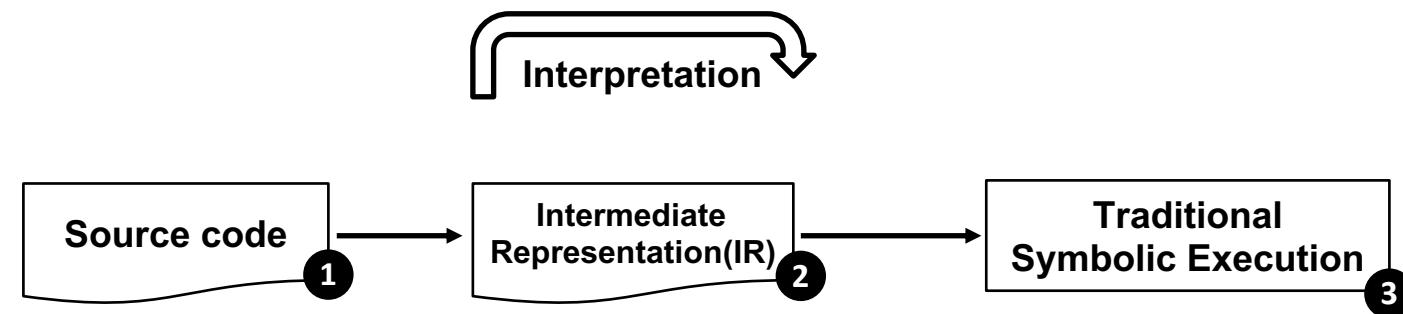
Examples of System & Application Software



- Q: How to improve the quality of software?
- A: Software testing



(Symbolic Execution)



A General Workflow of Symbolic Execution (SE)

One of the major concerns: **Performance**

Motivation

Motivation

- **Related work**

Motivation

- **Related work**

- Compiler optimization-based
 - Existing compiler optimization for accelerating symbolic execution [6,7]
 - Customized optimization for fast verification [8]

- **Related work**

- Compiler optimization-based
 - Existing compiler optimization for accelerating symbolic execution [6,7]
 - Customized optimization for fast verification [8]
- Others
 - Reducing path exploration and speeding up constraint solving

- **Related work**
 - Compiler optimization-based
 - Existing compiler optimization for accelerating symbolic execution [6,7]
 - Customized optimization for fast verification [8]
 - Others
 - Reducing path exploration and speeding up constraint solving
 - Limitations: they did not consider that interpreting alone can have performance downsides

- **Related work**
 - Compiler optimization-based
 - Existing compiler optimization for accelerating symbolic execution [6,7]
 - Customized optimization for fast verification [8]
 - Others
 - Reducing path exploration and speeding up constraint solving
 - Limitations: they did not consider that interpreting alone can have performance downsides
- **Summary**
 - Performance issues still remain in existing interpreting-based SE engines
 - Very few of them take the interpreting overhead alone into account

- **Related work**
 - Compiler optimization-based
 - Existing compiler optimization for accelerating symbolic execution [6,7]
 - Customized optimization for fast verification [8]
 - Others
 - Reducing path exploration and speeding up constraint solving
 - Limitations: they did not consider that interpreting alone can have performance downsides
- **Summary**
 - Performance issues still remain in existing interpreting-based SE engines
 - Very few of them take the interpreting overhead alone into account

How can we reduce interpretation overheads for fast symbolic execution?

Solution – FastKLEE (1/2)

Solution – FastKLEE (1/2)

- Two key insights

Solution – FastKLEE (1/2)

- **Two key insights**
 - **Insight 1:** The number of interpreted instructions tends to be stupendous

Solution – FastKLEE (1/2)

- Two key insights

- **Insight 1:** The number of interpreted instructions tends to be stupendous
 - several billion only in one hour run

```
Elapsed: 01: 00: 04
KLEE: done: explored paths = 125017
KLEE: done: avg. constructs per query = 74
KLEE: done: total queries = 8859
KLEE: done: valid queries = 6226
KLEE: done: invalid queries = 2633
KLEE: done: query cex = 8859

KLEE: done: total instructions = 605113213
KLEE: done: completed paths = 125017
KLEE: done: generated tests = 65
```

Solution – FastKLEE (1/2)

- Two key insights

- **Insight 1:** The number of interpreted instructions tends to be stupendous
 - several billion only in one hour run
- **Insight 2:** Only a small portion of pointers need bound checking during execution

```
Elapsed: 01: 00: 04
KLEE: done: explored paths = 125017
KLEE: done: avg. constructs per query = 74
KLEE: done: total queries = 8859
KLEE: done: valid queries = 6226
KLEE: done: invalid queries = 2633
KLEE: done: query cex = 8859

KLEE: done: total instructions = 605113213
KLEE: done: completed paths = 125017
KLEE: done: generated tests = 65
```

Solution – FastKLEE (1/2)

- Two key insights

- **Insight 1:** The number of interpreted instructions tends to be stupendous
 - several billion only in one hour run
- **Insight 2:** Only a small portion of pointers need bound checking during execution
 - Reduce interpreting overhead of most frequently interpreted ones (i.e., load/store instruction)

```
Elapsed: 01: 00: 04
KLEE: done: explored paths = 125017
KLEE: done: avg. constructs per query = 74
KLEE: done: total queries = 8859
KLEE: done: valid queries = 6226
KLEE: done: invalid queries = 2633
KLEE: done: query cex = 8859

KLEE: done: total instructions = 605113213
KLEE: done: completed paths = 125017
KLEE: done: generated tests = 65
```

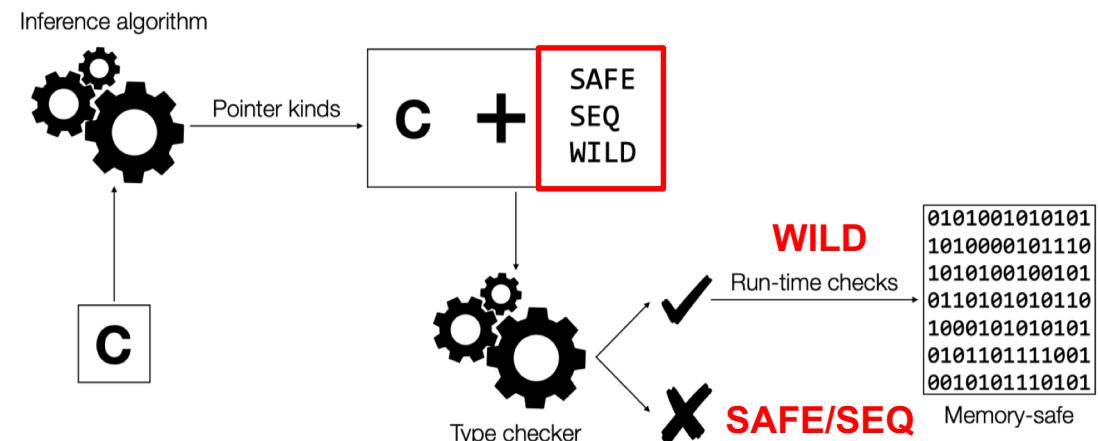
Solution – FastKLEE (1/2)

- Two key insights

- **Insight 1:** The number of interpreted instructions tends to be stupendous
 - several billion only in one hour run
- **Insight 2:** Only a small portion of pointers need bound checking during execution
 - Reduce interpreting overhead of most frequently interpreted ones (i.e., load/store instruction)
 - Inspired by **Type Inference** system

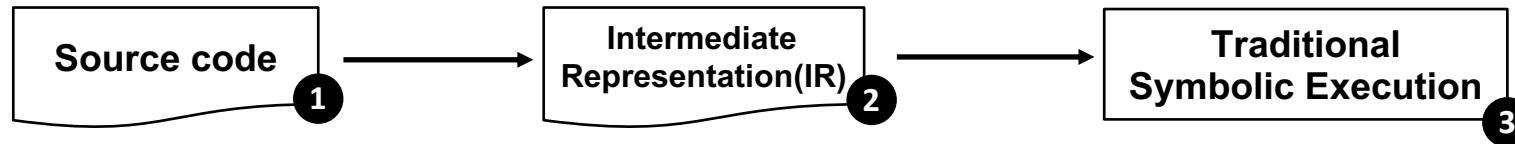
```
Elapsed: 01: 00: 04
KLEE: done: explored paths = 125017
KLEE: done: avg. constructs per query = 74
KLEE: done: total queries = 8859
KLEE: done: valid queries = 6226
KLEE: done: invalid queries = 2633
KLEE: done: query cex = 8859

KLEE: done: total instructions = 605113213
KLEE: done: completed paths = 125017
KLEE: done: generated tests = 65
```

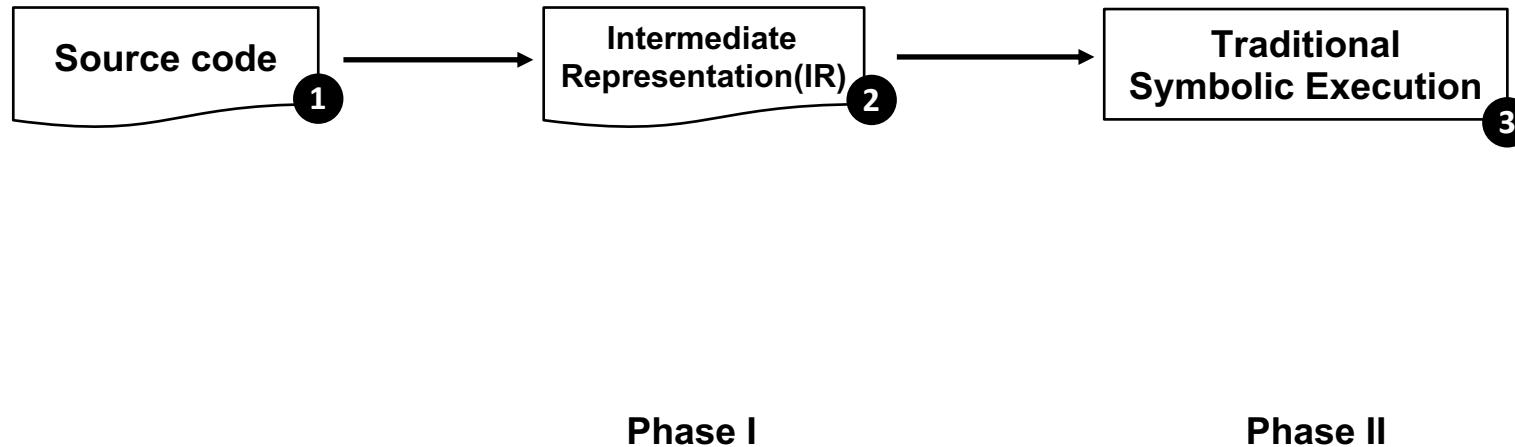


Solution – FastKLEE (2/2)

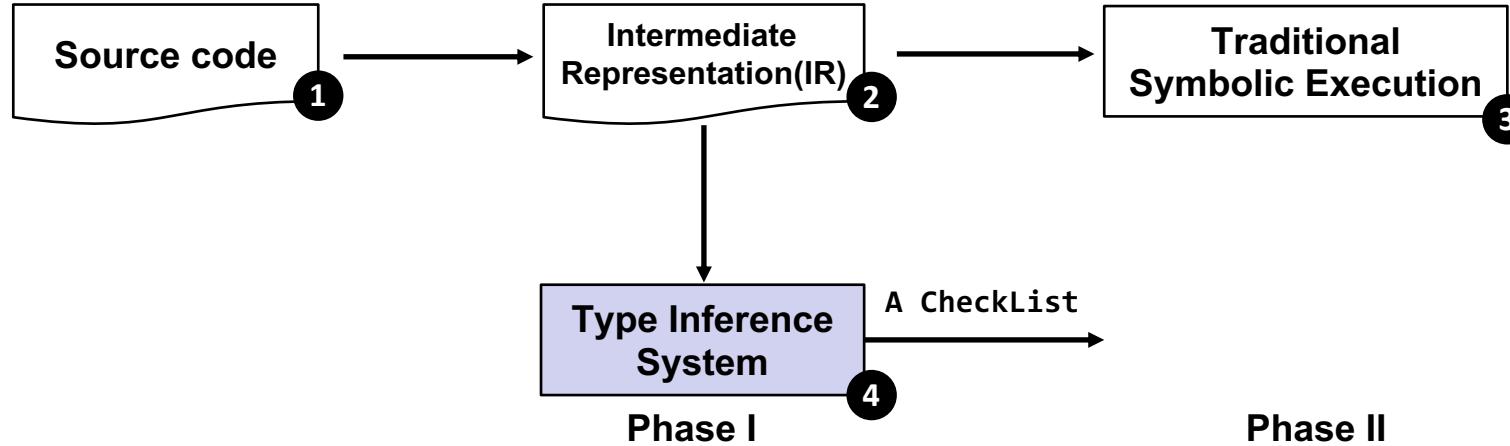
Solution – FastKLEE (2/2)



Solution – FastKLEE (2/2)

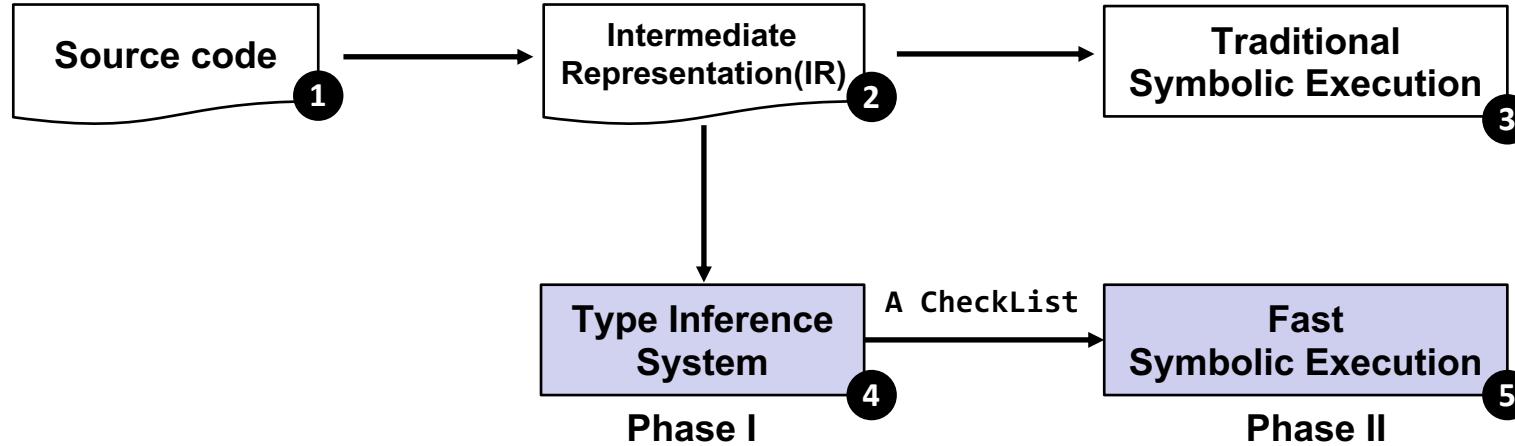


Solution – FastKLEE (2/2)



- ④ • **Phase I:** Introduce a **Type Inference System** to classify instruction types
 - **Unsafe** instructions will be stored in **CheckList**

Solution – FastKLEE (2/2)



- ④ • **Phase I:** Introduce a **Type Inference System** to classify instruction types
 - **Unsafe** instructions will be stored in **CheckList**

- ⑤ • **Phase II:** Conduct **Customized Instruction Operation** in Fast SE
 - Only perform checking for **Unsafe** instructions during interpretation

Preliminary Evaluation

Preliminary Evaluation

- **Implementation**
 - KLEE [1] and Ccured [4]
- **Benchmark**
 - GNU Coreutils

Preliminary Evaluation

- **Implementation**
 - KLEE [1] and Ccured [4]
- **Benchmark**
 - GNU Coreutils
- **Metric**
 - **Speedups**: the **time** spent on exploring the same number of instructions

$$\text{Speedups} : \frac{T_{baseline} - T_{our}}{T_{baseline}} \times 100$$

- $T_{baseline}$: **existing approach**
- T_{our} : **our approach**

Preliminary Evaluation

- **Implementation**

- KLEE [1] and Ccured [4]

- **Benchmark**

- GNU Coreutils

- **Metric**

- **Speedups**: the time spent on exploring the same number of instructions

$$\text{Speedups} : \frac{T_{baseline} - T_{our}}{T_{baseline}} \times 100$$

- $T_{baseline}$: existing approach
- T_{our} : our approach

- **Results**

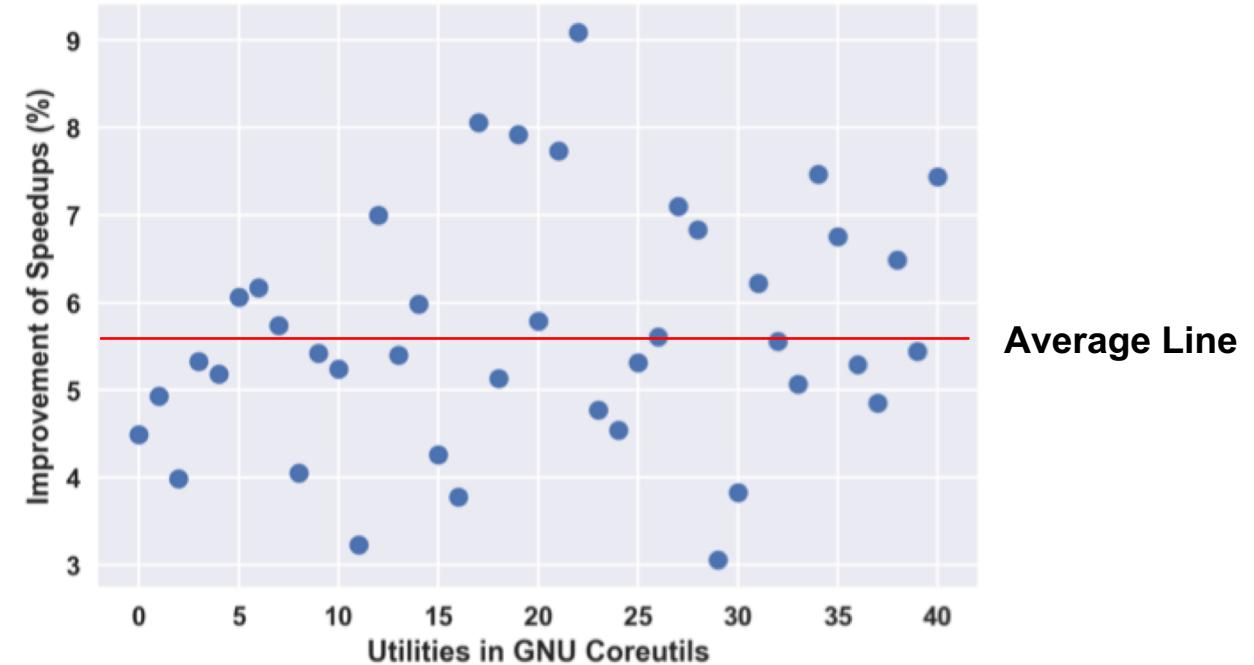


Figure 2: Scatter plot of the improvement in speedups

Preliminary Evaluation

- **Implementation**

- KLEE [1] and Ccured [4]

- **Benchmark**

- GNU Coreutils

- **Metric**

- **Speedups**: the time spent on exploring the same number of instructions

$$\text{Speedups} : \frac{T_{baseline} - T_{our}}{T_{baseline}} \times 100$$

- $T_{baseline}$: existing approach
- T_{our} : our approach

- **Results**

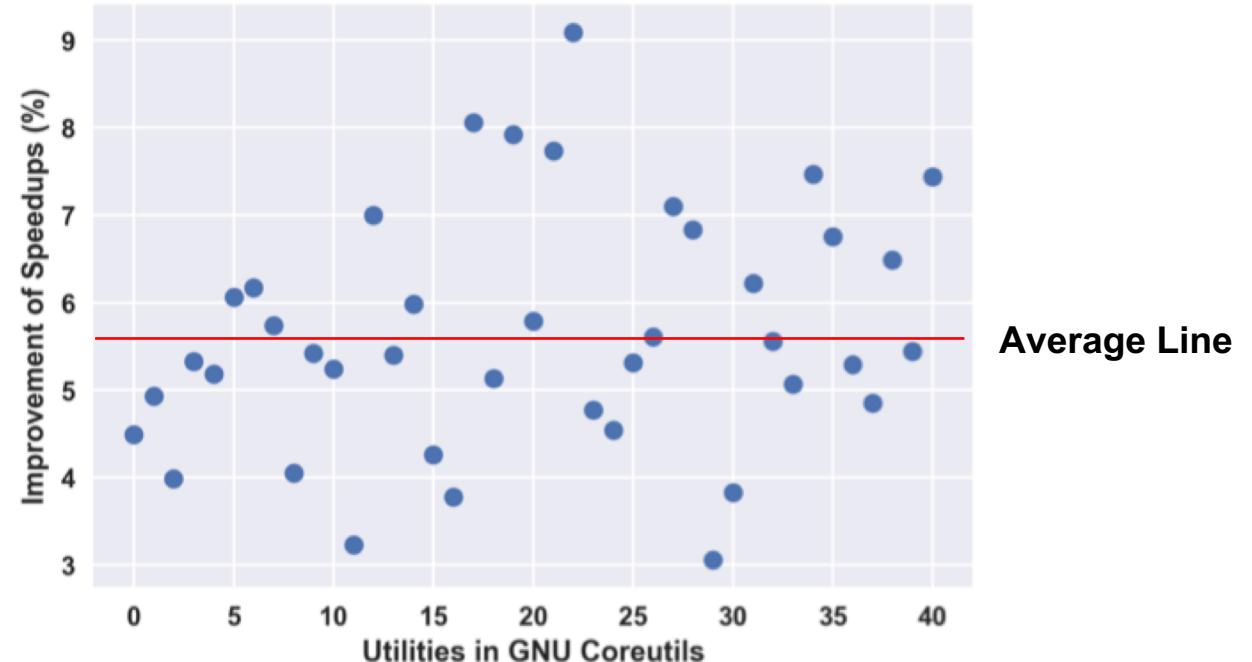


Figure 2: Scatter plot of the improvement in speedups

- FastKLEE is able to reduce by up to 9.1% (5.6% on average) time as the state-of-the-art approach

Conclusion

• Contribution

- We present **FastKLEE**, which aims to reduce the interpretation overheads for fast symbolic execution

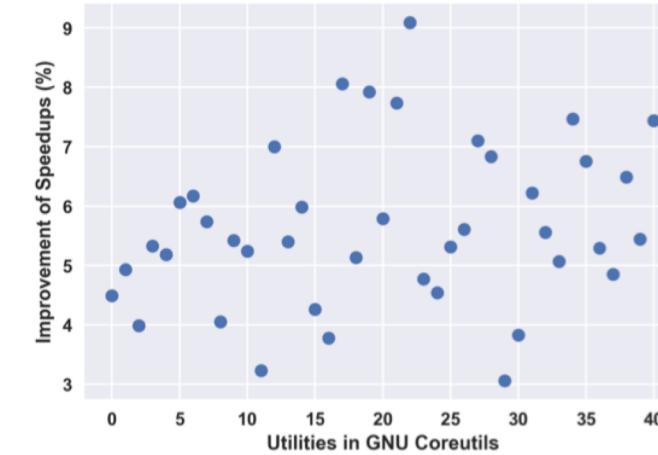
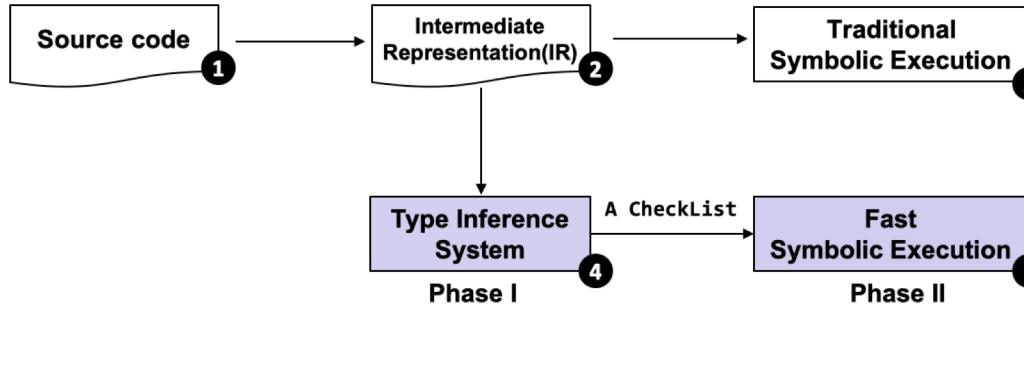


Figure 2: Scatter plot of the improvement in speedups

• Future work

- Use FastKLEE to explore more **valuable** execution paths in software systems
 - valuable: vulnerable and exploitable

Code: <https://github.com/haoxitu/FastKLEE>

Email: haoxitu.2020@phdcs.smu.edu.sg

(Please feel free to pull requests or raise any questions if you have!)

References

- [1] Cristian Cadar, Daniel Dunbar, and Dawson Engler. 2008. **KLEE: unassisted and automatic generation of high-coverage tests for complex systems programs.** In Proceedings of the 8th USENIX conference on Operating systems design and implementation (OSDI'08). USENIX Association, USA, 209–224.
- [2] V. Chipounov, V. Kuznetsov, and G. Cadea, **S2E: a platform for in-vivo multi-path analysis of software systems**, in Proceedings of the sixteenth international conference on Architectural support for programming languages and operating systems, New York, NY, USA, Mar. 2011, pp. 265–278.
- [3] Y. Shoshtaishvili et al., **SOK: (State of) The Art of War: Offensive Techniques in Binary Analysis**, in 2016 IEEE Symposium on Security and Privacy (SP), May 2016, pp. 138–157.
- [4] Ccured: George C. Necula, Jeremy Condit, Matthew Harren, Scott McPeak, and Westley Weimer. 2005. **CCured: type-safe retrofitting of legacy software**. ACM Trans. Program. Lang. Syst. 27, 3 (May 2005), 477–526.
- [5] S. Poeplau and A. Francillon, **Symbolic execution with SymCC: Don't interpret, compile!**, in 29th USENIX Security Symposium, 2020, pp. 181–198.
- [6] Dong, Shiyu, Oswaldo Olivo, Lingming Zhang, and Sarfraz Khurshid. **Studying the influence of standard compiler optimizations on symbolic execution**, In 2015 IEEE 26th International Symposium on Software Reliability Engineering (ISSRE). IEEE, 2015.
- [7] Chen, Junjie, Wenxiang Hu, Lingming Zhang, Dan Hao, Sarfraz Khurshid, and Lu Zhang. **Learning to accelerate symbolic execution via code transformation**, In 32nd European Conference on Object-Oriented Programming (ECOOP 2018).
- [8] Wagner, Jonas, Volodymyr Kuznetsov, and George Cadea. **{-VERIFY}: Optimizing Programs for Fast {Verification}**, In 14th Workshop on Hot Topics in Operating Systems (HotOS XIV). 2013.

Thanks you!

FastKLEE: Faster Symbolic Execution via Reducing Redundant Bound Checking of Type-Safe Pointers

Haoxin Tu, Lingxiao Jiang, Xuhua Ding (Singapore Management University)
He Jiang (Dalian University of Technology)

Code: <https://github.com/haoxitu/FastKLEE>

Email: haoxitu.2020@phdcs.smu.edu.sg

(Please feel free to pull requests or raise any questions if you have!)

Backup: other evaluations

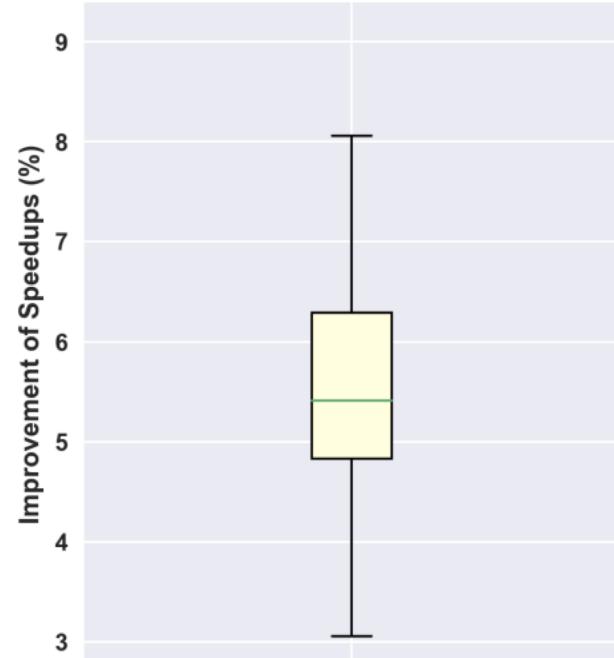


Figure 4: Box plot of the improvement in speedups

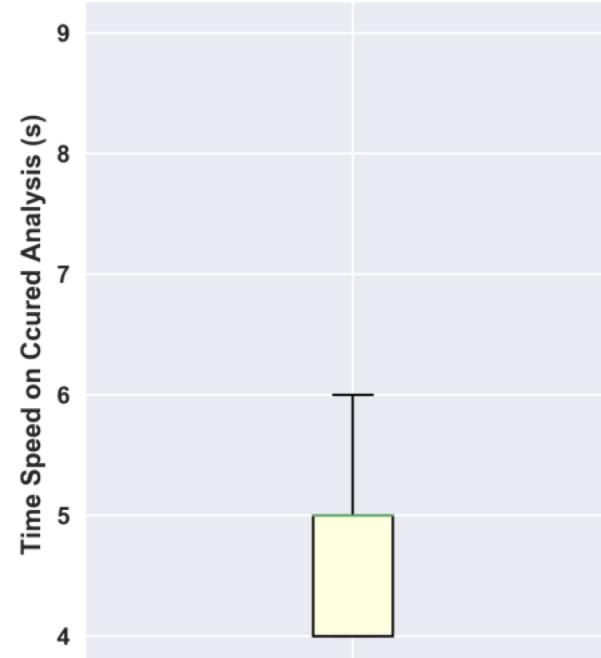
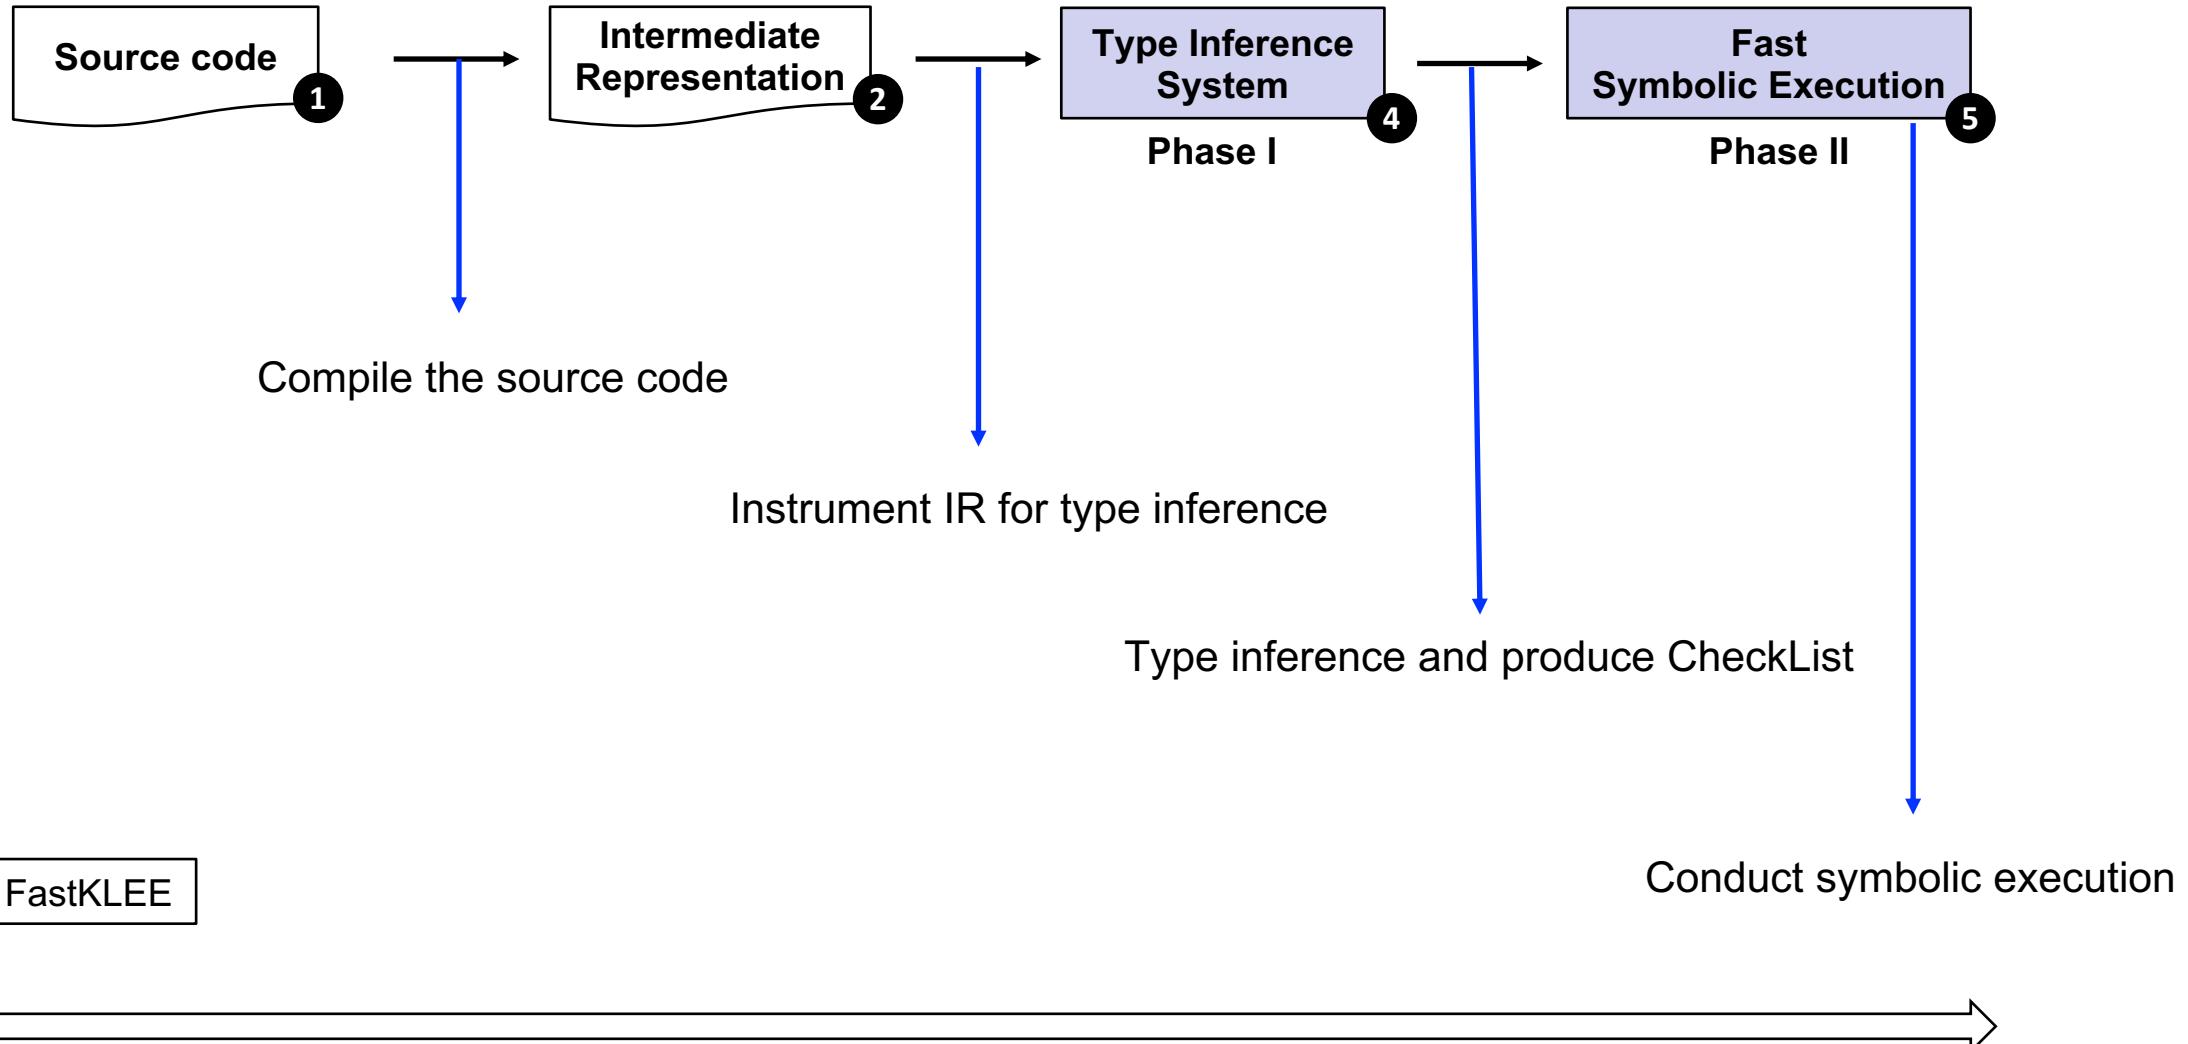


Figure 5: Box plot of the time spent on type inference

Backup: Tool Usage

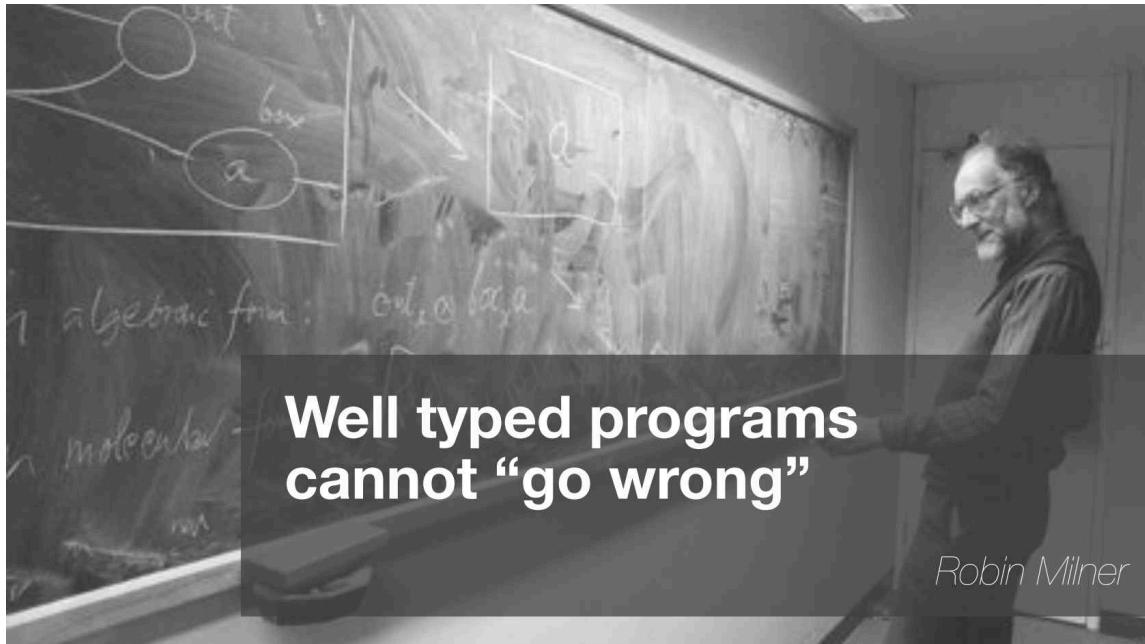


Code: <https://github.com/haoxitu/FastKLEE>

(Please feel free to pull requests and raise any questions if you have!)

Background: type safety (1)

- What is type safety?
 - Type safety means that the compiler will validate types while compiling, and throw an error if assign the wrong type to a variable.



Source: [What is Type-safety?](#)

 Java

```
int x = 100;
String y = "hello";
int z = x + y;
```

Error: incompatible types:
String cannot be converted to int

Type **safe** programming language

 C

```
int x = 100;
char* y = "hello";
int z = x + y;
```

No errors

Type **unsafe** programming language

Note: type safety was intentionally sacrificed for flexibility/performance in C

Backup: type inference (2)

- **Type inference**
 - General Definition: automatic detection of the type of an expression in a programming language
 - Type inference in C
 - Try to keep type safety (specifically for pointer types)
- **CCured [4]**
 - CCured transforms C programs to achieve memory safety guarantees
 - A large portion of pointers can be statically verified to be safe, and only small need run-time check

