

FastKLEE: Faster Symbolic Execution via Redundant Bound Checking of Type-Safe Pointers

ESEC/FSE 2022 Tool Demonstration

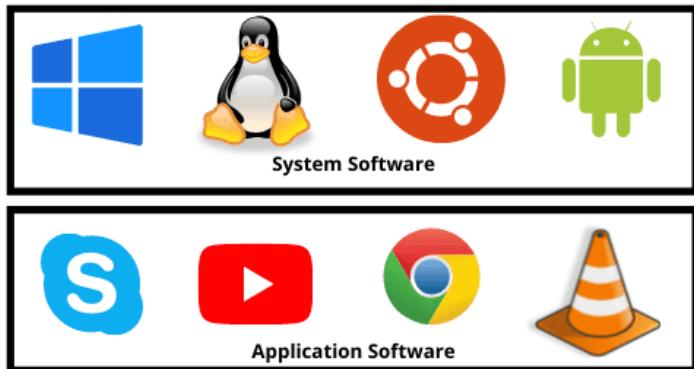
Haoxin Tu, Lingxiao Jiang, Xuhua Ding (Singapore Management University)
He Jiang (Dalian University of Technology)



Background

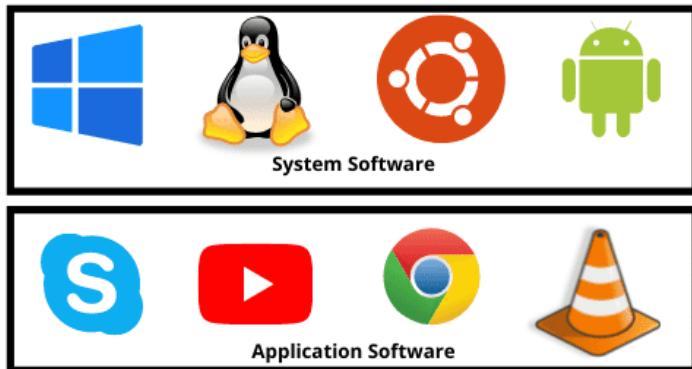
Background

Examples of System & Application Software



Background

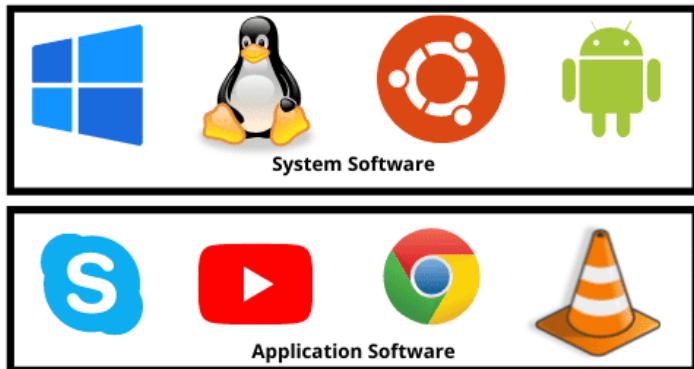
Examples of System & Application Software



- Q: How to improve the quality of software?

Background

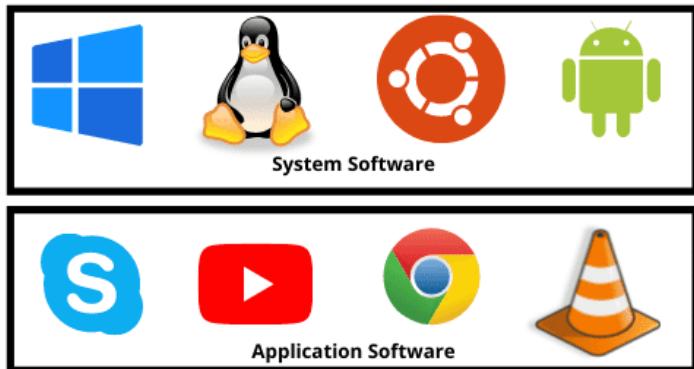
Examples of System & Application Software



- Q: How to improve the quality of software?
- A: Software testing

Background

Examples of System & Application Software



- Q: How to improve the quality of software?
- A: Software testing



(Symbolic Execution)

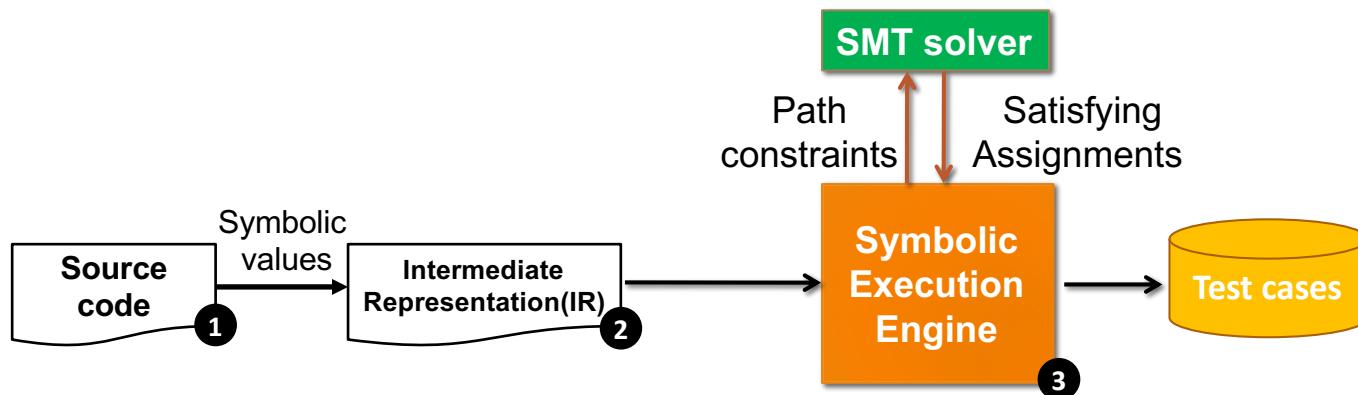
Background

Examples of System & Application Software



- Q: How to improve the quality of software?
- A: Software testing

↓
(Symbolic Execution)



- General workflow of traditional symbolic execution engine (e.g., KLEE)

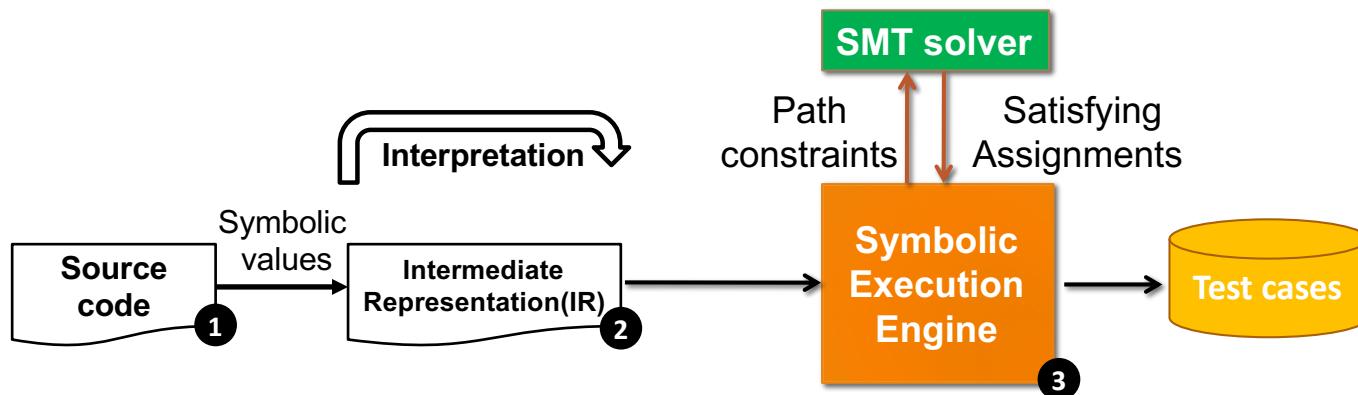
Background

Examples of System & Application Software



- Q: How to improve the quality of software?
- A: Software testing

↓
(Symbolic Execution)



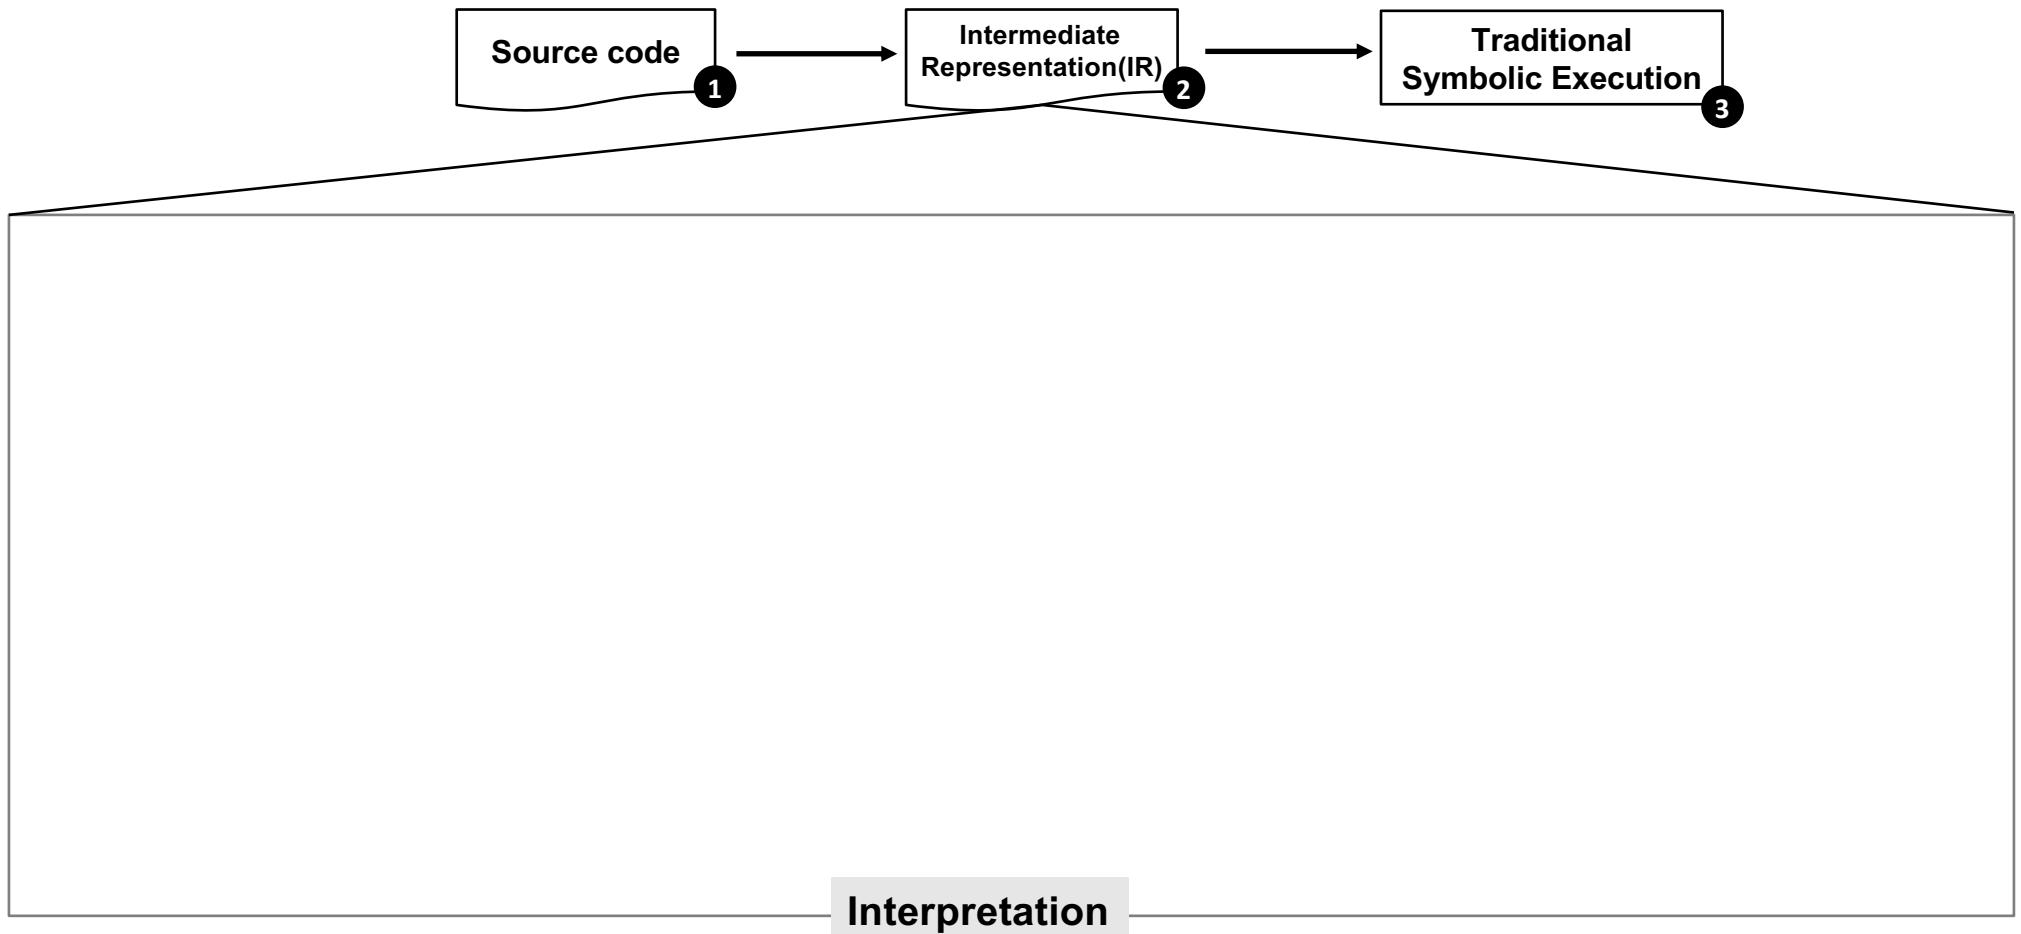
- General workflow of traditional symbolic execution engine (e.g., KLEE)

Motivation

Motivation



Motivation



Motivation



(1) Observation

- The number of interpreted instructions tends to be **huge** (several billion only in one hour run)

```
Elapsed: 01:00:04
KLEE: done: explored paths = 125017
KLEE: done: avg. constructs per query = 74
KLEE: done: total queries = 8859
KLEE: done: valid queries = 6226
KLEE: done: invalid queries = 2633
KLEE: done: query cex = 8859

KLEE: done: total instructions = 605113213
KLEE: done: completed paths = 125017
KLEE: done: generated tests = 65
```

Interpretation

Motivation



(1) Observation

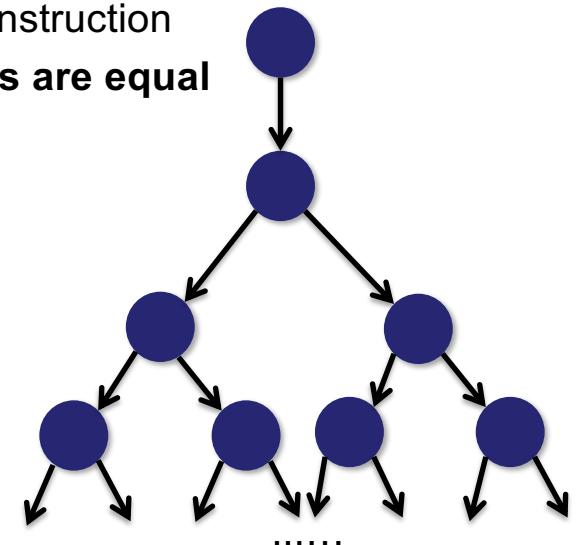
- The number of interpreted instructions tends to be **huge** (several billion only in one hour run)

```
Elapsed: 01:00:04
KLEE: done: explored paths = 125017
KLEE: done: avg. constructs per query = 74
KLEE: done: total queries = 8859
KLEE: done: valid queries = 6226
KLEE: done: invalid queries = 2633
KLEE: done: query cex = 8859

KLEE: done: total instructions = 605113213
KLEE: done: completed paths = 125017
KLEE: done: generated tests = 65
```

(2) Overheads in current symbolic execution

- The color depth represents the overheads of an interpreted instruction
- **All instructions are equal**



Interpretation

Motivation



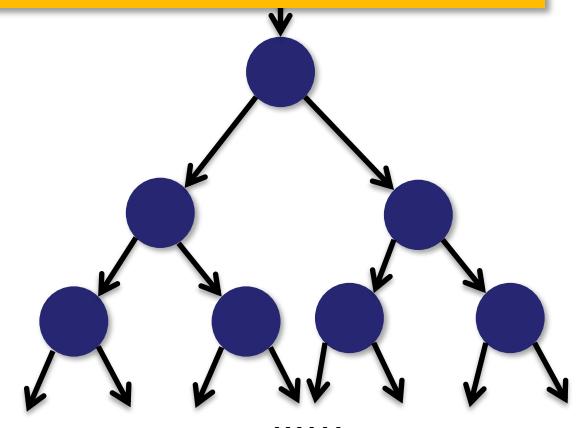
(1) Observation

(2) Overheads in current symbolic execution

Can we reduce the overheads of interpreted instructions for faster symbolic execution?

```
Elapsed: 01:00:04
KLEE: done: explored paths = 125017
KLEE: done: avg. constructs per query = 74
KLEE: done: total queries = 8859
KLEE: done: valid queries = 6226
KLEE: done: invalid queries = 2633
KLEE: done: query cex = 8859

KLEE: done: total instructions = 605113213
KLEE: done: completed paths = 125017
KLEE: done: generated tests = 65
```



Interpretation

Solution – FastKLEE (1/2)

[1] Cured: George C. Necula, Jeremy Condit, Matthew Harren, Scott McPeak, and Westley Weimer. 2005. **CCured: type-safe retrofitting of legacy software**. ACM Trans. Program. Lang. Syst. 27, 3 (May 2005), 477–526.

Solution – FastKLEE (1/2)

- Key insight

[1] Cured: George C. Necula, Jeremy Condit, Matthew Harren, Scott McPeak, and Westley Weimer. 2005. **CCured: type-safe retrofitting of legacy software**. ACM Trans. Program. Lang. Syst. 27, 3 (May 2005), 477–526.

Solution – FastKLEE (1/2)

- **Key insight**
 - Only a small portion of memory-related instructions need bound checking

[1] Cured: George C. Necula, Jeremy Condit, Matthew Harren, Scott McPeak, and Westley Weimer. 2005. **CCured: type-safe retrofitting of legacy software**. ACM Trans. Program. Lang. Syst. 27, 3 (May 2005), 477–526.

Solution – FastKLEE (1/2)

- **Key insight**
 - Only a small portion of memory-related instructions need bound checking
 - Reduce interpreting overhead of most frequently interpreted ones (i.e., load/store instruction)

[1] Cured: George C. Necula, Jeremy Condit, Matthew Harren, Scott McPeak, and Westley Weimer. 2005. **CCured: type-safe retrofitting of legacy software**. ACM Trans. Program. Lang. Syst. 27, 3 (May 2005), 477–526.

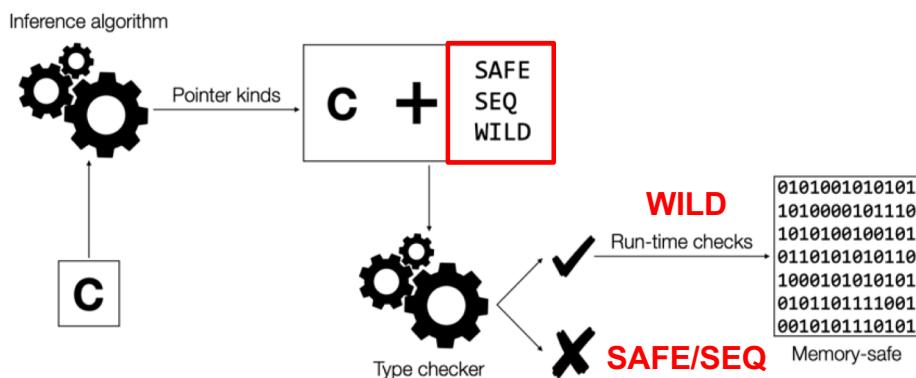
Solution – FastKLEE (1/2)

- **Key insight**
 - Only a small portion of memory-related instructions need bound checking
 - Reduce interpreting overhead of most frequently interpreted ones (i.e., load/store instruction)
 - Inspired by **Type Inference** system [1]

[1] Cured: George C. Necula, Jeremy Condit, Matthew Harren, Scott McPeak, and Westley Weimer. 2005. CCured: type-safe retrofitting of legacy software. ACM Trans. Program. Lang. Syst. 27, 3 (May 2005), 477–526.

Solution – FastKLEE (1/2)

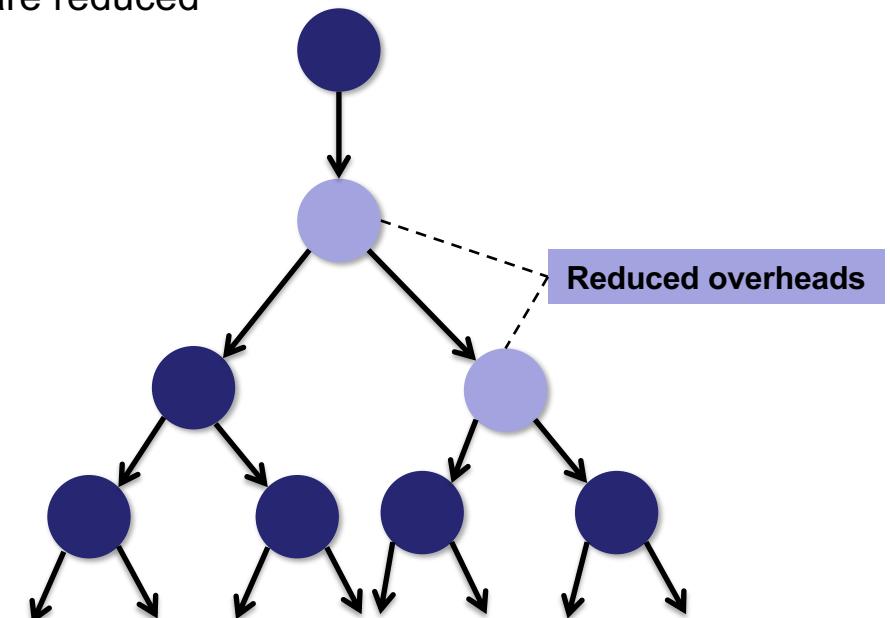
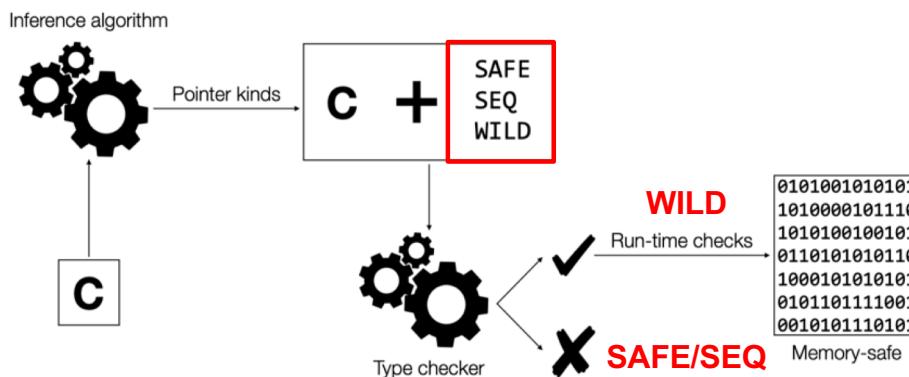
- Key insight
 - Only a small portion of memory-related instructions need bound checking
 - Reduce interpreting overhead of most frequently interpreted ones (i.e., load/store instruction)
 - Inspired by **Type Inference** system [1]



[1] Ccured: George C. Necula, Jeremy Condit, Matthew Harren, Scott McPeak, and Westley Weimer. 2005. CCured: type-safe retrofitting of legacy software. ACM Trans. Program. Lang. Syst. 27, 3 (May 2005), 477–526.

Solution – FastKLEE (1/2)

- **Key insight**
 - Only a small portion of memory-related instructions need bound checking
 - Reduce interpreting overhead of most frequently interpreted ones (i.e., load/store instruction)
 - Inspired by **Type Inference** system [1]
- **Advantage: overheads in FastKLEE**
 - Interpretation overheads for some instructions are reduced



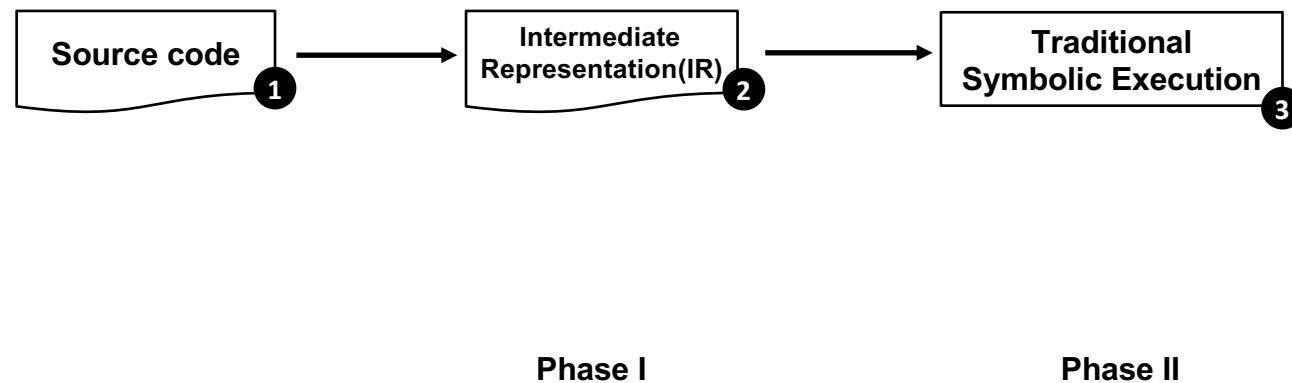
[1] CCured: George C. Necula, Jeremy Condit, Matthew Harren, Scott McPeak, and Westley Weimer. 2005. CCured: type-safe retrofitting of legacy software. ACM Trans. Program. Lang. Syst. 27, 3 (May 2005), 477–526.

Solution – FastKLEE (2/2)

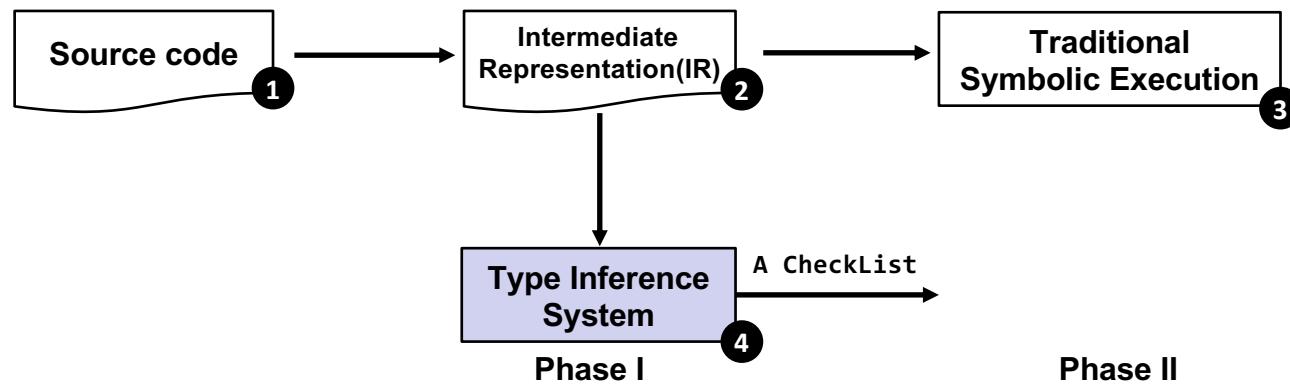
Solution – FastKLEE (2/2)



Solution – FastKLEE (2/2)

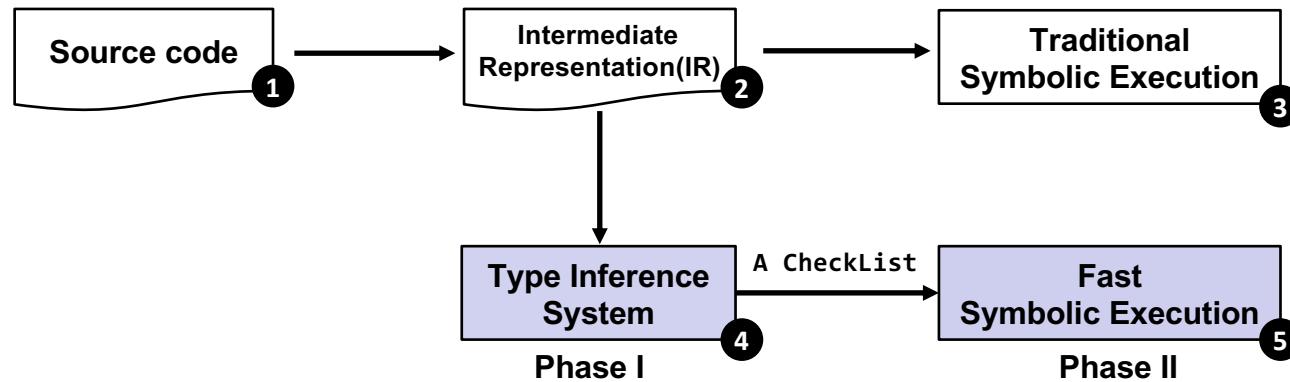


Solution – FastKLEE (2/2)



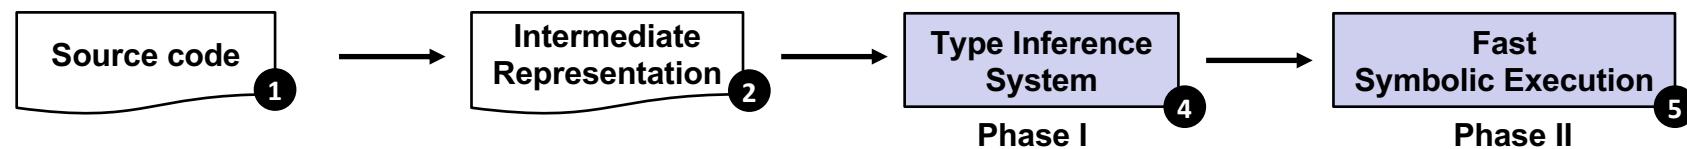
- ④ • **Phase I:** Introduce a **Type Inference System** to classify memory-related instruction types
 - **Unsafe** memory instructions will be stored in **CheckList**

Solution – FastKLEE (2/2)



- ④ • **Phase I:** Introduce a **Type Inference System** to classify memory-related instruction types
 - **Unsafe** memory instructions will be stored in **CheckList**
- ⑤ • **Phase II:** Conduct **Customized Memory Operation** in Fast symbolic execution
 - Only perform checking for **Unsafe** memory instructions during interpretation

Tool demonstration



- **Implementation**
 - KLEE [1] and Ccured [2]

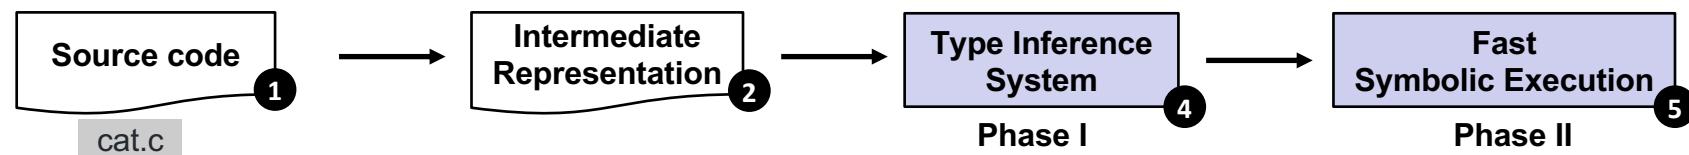
Installation of FastKLEE

Full video demo: https://youtu.be/fjV_a3kt-mo

[1] Cristian Cadar, Daniel Dunbar, and Dawson Engler. 2008. **KLEE: unassisted and automatic generation of high-coverage tests for complex systems programs**. In OSDI'08. USENIX Association, USA, 209–224.

[2] George C. Necula, Jeremy Condit, Matthew Harren, Scott McPeak, and Westley Weimer. 2005. **CCured: type-safe retrofitting of legacy software**. ACM Trans. Program. Lang. Syst. 27, 3 (May 2005), 477–526.

Tool demonstration



- **Implementation**
 - KLEE [1] and Ccured [2]

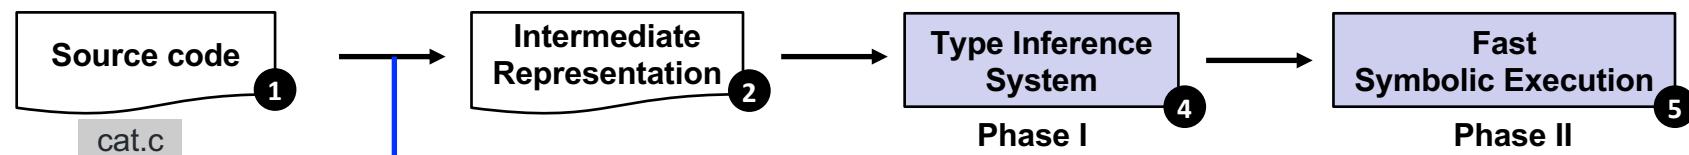
Installation of FastKLEE

Full video demo: https://youtu.be/fjV_a3kt-mo

[1] Cristian Cadar, Daniel Dunbar, and Dawson Engler. 2008. **KLEE: unassisted and automatic generation of high-coverage tests for complex systems programs**. In OSDI'08. USENIX Association, USA, 209–224.

[2] George C. Necula, Jeremy Condit, Matthew Harren, Scott McPeak, and Westley Weimer. 2005. **CCured: type-safe retrofitting of legacy software**. ACM Trans. Program. Lang. Syst. 27, 3 (May 2005), 477–526.

Tool demonstration



1. Compile the source code

```
$ clang -emit-llvm -c cat.c -o cat.bc
```

- Implementation

- KLEE [1] and Ccured [2]

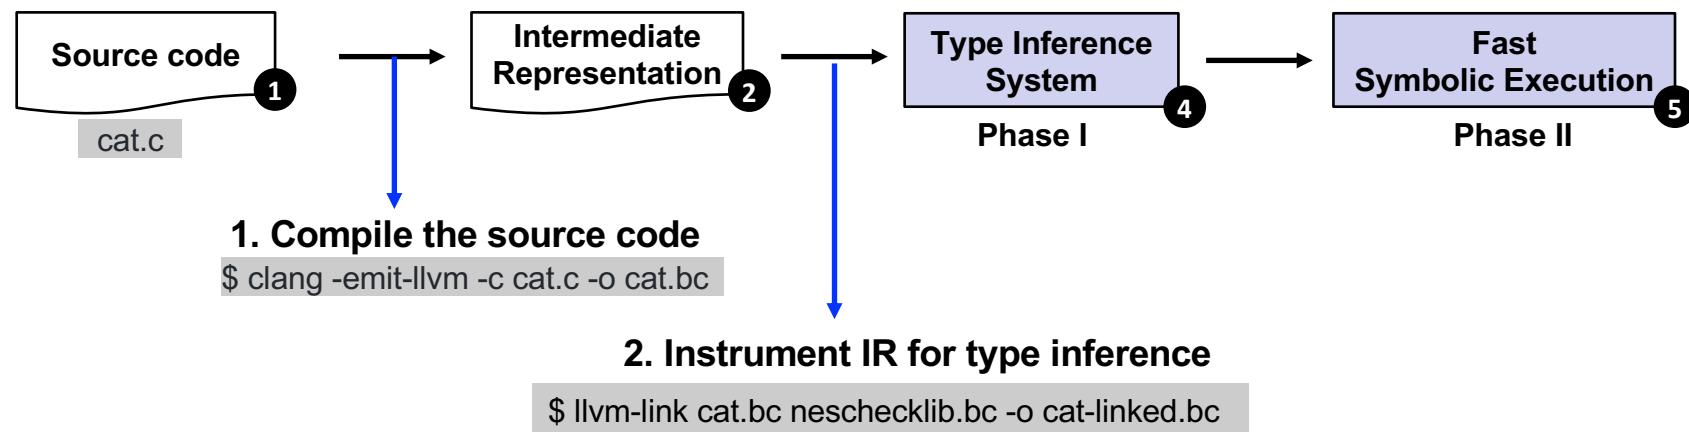
Installation of FastKLEE

Full video demo: https://youtu.be/fjV_a3kt-mo

[1] Cristian Cadar, Daniel Dunbar, and Dawson Engler. 2008. **KLEE: unassisted and automatic generation of high-coverage tests for complex systems programs**. In OSDI'08. USENIX Association, USA, 209–224.

[2] Ccured: George C. Necula, Jeremy Condit, Matthew Harren, Scott McPeak, and Westley Weimer. 2005. **CCured: type-safe retrofitting of legacy software**. ACM Trans. Program. Lang. Syst. 27, 3 (May 2005), 477–526.

Tool demonstration



- **Implementation**

- KLEE [1] and Ccured [2]

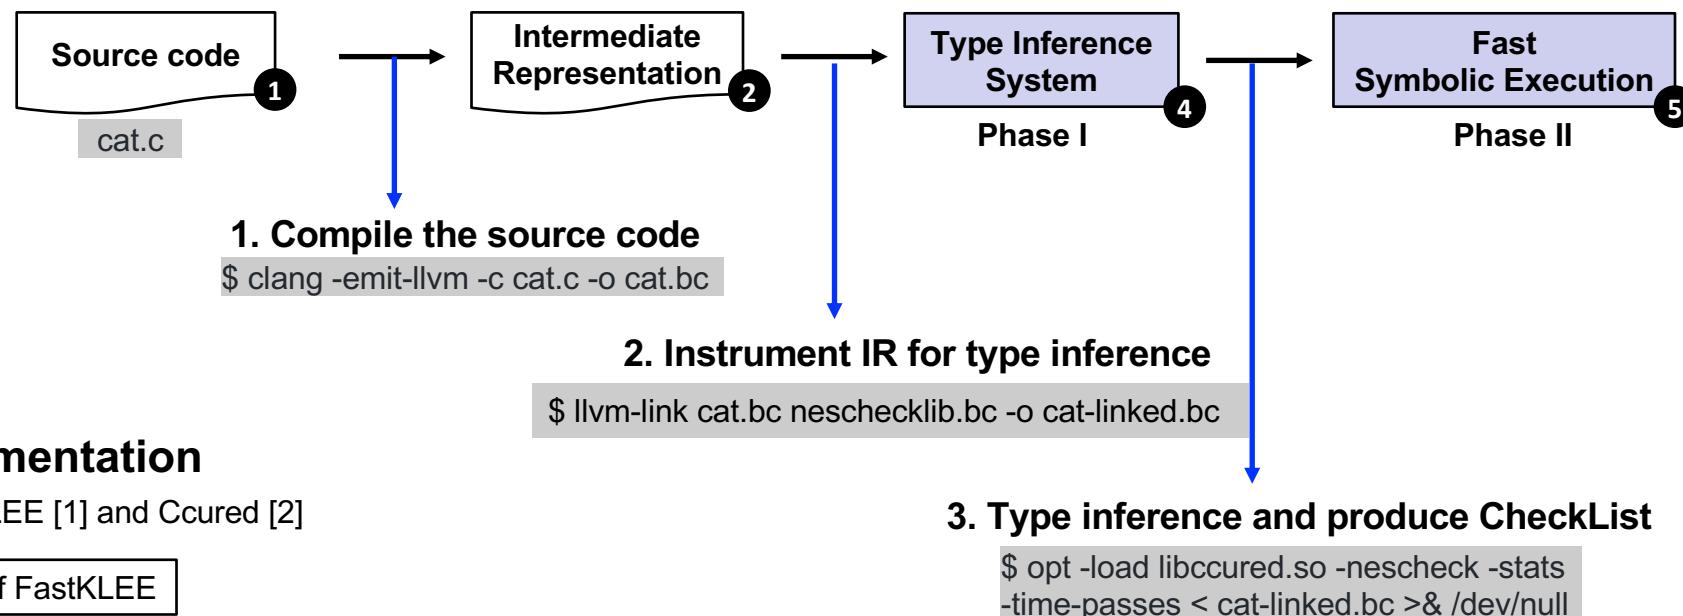
Installation of FastKLEE

Full video demo: https://youtu.be/fjV_a3kt-mo

[1] Cristian Cadar, Daniel Dunbar, and Dawson Engler. 2008. **KLEE: unassisted and automatic generation of high-coverage tests for complex systems programs**. In OSDI'08. USENIX Association, USA, 209–224.

[2] Ccured: George C. Necula, Jeremy Condit, Matthew Harren, Scott McPeak, and Westley Weimer. 2005. **CCured: type-safe retrofitting of legacy software**. ACM Trans. Program. Lang. Syst. 27, 3 (May 2005), 477–526.

Tool demonstration



- **Implementation**

- KLEE [1] and Ccured [2]

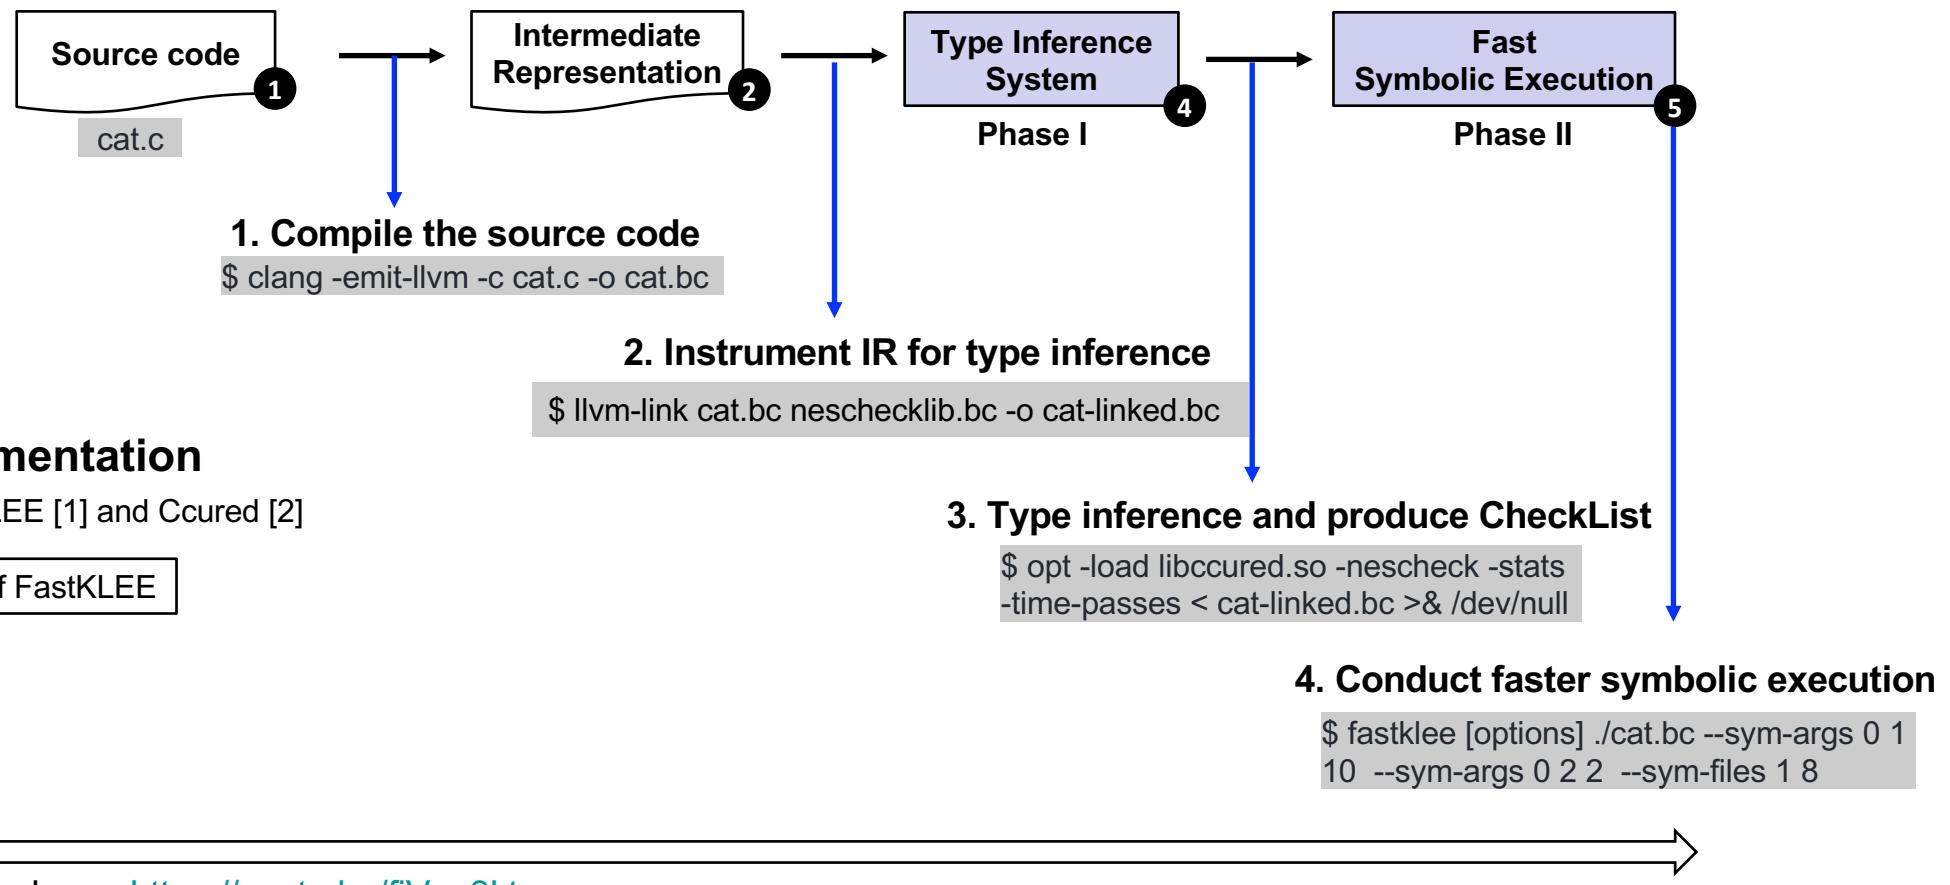
Installation of FastKLEE

Full video demo: https://youtu.be/fjV_a3kt-mo

[1] Cristian Cadar, Daniel Dunbar, and Dawson Engler. 2008. **KLEE: unassisted and automatic generation of high-coverage tests for complex systems programs**. In OSDI'08. USENIX Association, USA, 209–224.

[2] Ccured: George C. Necula, Jeremy Condit, Matthew Harren, Scott McPeak, and Westley Weimer. 2005. **CCured: type-safe retrofitting of legacy software**. ACM Trans. Program. Lang. Syst. 27, 3 (May 2005), 477–526.

Tool demonstration



Full video demo: https://youtu.be/fjV_a3kt-mo

[1] Cristian Cadar, Daniel Dunbar, and Dawson Engler. 2008. **KLEE: unassisted and automatic generation of high-coverage tests for complex systems programs**. In OSDI'08. USENIX Association, USA, 209–224.

[2] Ccured: George C. Necula, Jeremy Condit, Matthew Harren, Scott McPeak, and Westley Weimer. 2005. **CCured: type-safe retrofitting of legacy software**. ACM Trans. Program. Lang. Syst. 27, 3 (May 2005), 477–526.

Preliminary Evaluation

Preliminary Evaluation

- **Benchmark**

- GNU Coreutils
- ~ 1-5k SLOC for each test program

Preliminary Evaluation

- **Benchmark**

- GNU Coreutils
- ~ 1-5k SLOC for each test program

- **Metric**

- **Speedups**: the **time** spent on exploring the same number of instructions

$$\text{Speedups} : \frac{T_{baseline} - T_{our}}{T_{baseline}} \times 100$$

- $T_{baseline}$: **existing approach**
- T_{our} : **our approach**

Preliminary Evaluation

- **Benchmark**

- GNU Coreutils
- ~ 1-5k SLOC for each test program

- **Metric**

- **Speedups**: the **time** spent on exploring the same number of instructions

$$\text{Speedups} : \frac{T_{baseline} - T_{our}}{T_{baseline}} \times 100$$

- $T_{baseline}$: existing approach
- T_{our} : our approach

- **Results**

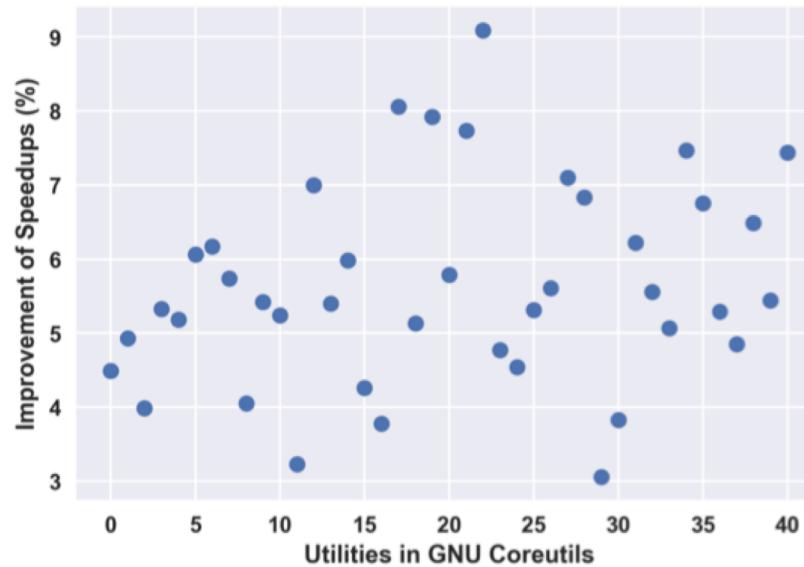


Figure 2: Scatter plot of the improvement in speedups

Preliminary Evaluation

- **Benchmark**

- GNU Coreutils
- ~ 1-5k SLOC for each test program

- **Metric**

- **Speedups**: the **time** spent on exploring the same number of instructions

$$\text{Speedups} : \frac{T_{baseline} - T_{our}}{T_{baseline}} \times 100$$

- $T_{baseline}$: existing approach
- T_{our} : our approach

- **Results**

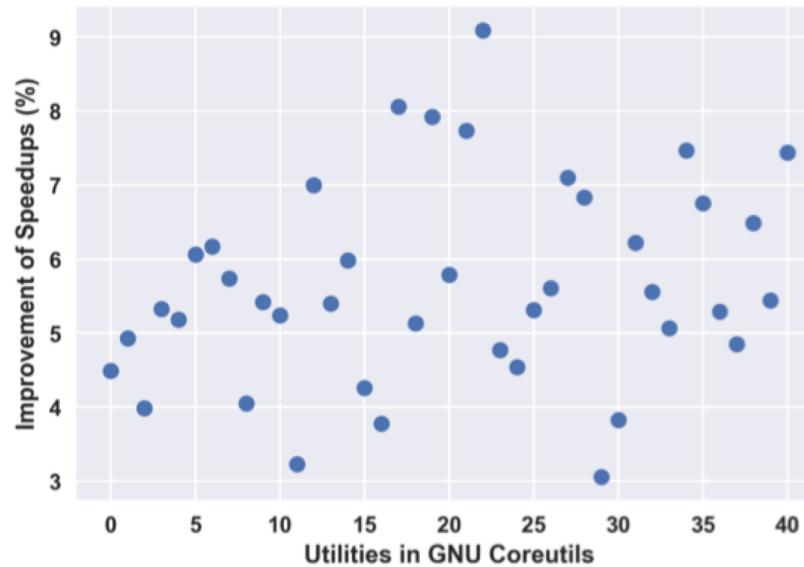


Figure 2: Scatter plot of the improvement in speedups

- FastKLEE can reduce by up to **9.1%** time as the state-of-the-art approach (i.e., KLEE)

Conclusion

- **Contribution**

- We present **FastKLEE**, which reduces the interpretation overheads for faster symbolic execution

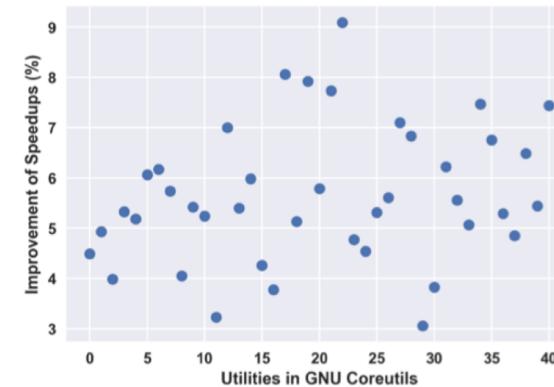
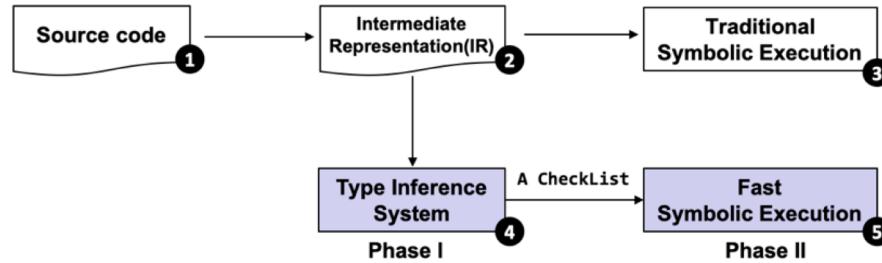


Figure 2: Scatter plot of the improvement in speedups

- **Future work**

- Use FastKLEE to explore more **valuable** execution paths in software systems
 - valuable: vulnerable and exploitable

Code: <https://github.com/haoxitu/FastKLEE>

Video demo: https://youtu.be/fjV_a3kt-mo

Email: haoxitu.2020@phdcs.smu.edu.sg

(Please feel free to pull requests or raise any questions if you have!)